

1. Regresión logística

2. K-means



Meetup Monterrey Data Science & Engineering



CoWo, Monterrey, Nuevo León, México



22 de junio de 2017



Rafael Rodríguez Morales (rafarodrz@gmail.com)

Aprendizaje automático

	<u>Unsupervised</u>	<u>Supervised</u>
<u>Continuous</u>	<ul style="list-style-type: none">• Clustering & Dimensionality Reduction<ul style="list-style-type: none">○ SVD○ PCA○ K-means	<ul style="list-style-type: none">• Regression<ul style="list-style-type: none">○ Linear○ Polynomial• Decision Trees• Random Forests
<u>Categorical</u>	<ul style="list-style-type: none">• Association Analysis<ul style="list-style-type: none">○ Apriori○ FP-Growth• Hidden Markov Model	<ul style="list-style-type: none">• Classification<ul style="list-style-type: none">○ KNN○ Trees○ Logistic Regression○ Naive-Bayes○ SVM

1. Regresión logística

0. Regresión

Regresión

- Tenemos: training set

$$\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$$


$$X = \begin{bmatrix} \text{---} (x^{(1)})^T \text{---} \\ \text{---} (x^{(2)})^T \text{---} \\ \vdots \\ \text{---} (x^{(m)})^T \text{---} \end{bmatrix} \quad \vec{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix}$$

La naturaleza de y
define si es
regresión lineal o
regresión logística

- Queremos: los parámetros para la función

$$h_{\theta}(x)$$

Regresión



Regresión lineal

- Hipótesis


$$h(x) = \sum_{i=0}^n \theta_i x_i = \theta^T x$$

$$h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$

- Función costo (LMS)

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$


Regresión lineal



Minimizar la función costo


- Gradient descent
- BFGS
- L-BFGS
- Stochastic Average Gradient

Gradient descent idea general



Gradient descent para regresión lineal

- Iterativo
- Parte de vector θ inicial y hace actualizaciones


$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

Repeat until convergence {


$$\theta_j := \theta_j + \alpha \sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)} \quad (\text{for every } j)$$

}

Convergencia



1. Regresión logística



Regresión

- Tenemos: training set

$$\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$$

$$X = \begin{bmatrix} \text{---} (x^{(1)})^T \text{---} \\ \text{---} (x^{(2)})^T \text{---} \\ \vdots \\ \text{---} (x^{(m)})^T \text{---} \end{bmatrix} \quad \vec{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix}$$

La naturaleza de y
define si es
regresión lineal o
regresión logística

- Queremos: los parámetros para la función


$$h_{\theta}(x)$$

Regresión logística

- Hipótesis


$$h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$

$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$



Regresión logística

Umbral





$$P(y = 1 \mid x; \theta) = h_{\theta}(x)$$

$$P(y = 0 \mid x; \theta) = 1 - h_{\theta}(x)$$

Regresión logística

- Función costo


$$J(\theta) = -\frac{1}{m} \left(\sum_{i=1}^m y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right)$$



Gradient descent para regresión logística

- Iterativo
- Parte de vector θ inicial y hace actualizaciones

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$



$$\sum_{i=1}^m y^{(i)} \log h(x^{(i)}) + (1 - y^{(i)}) \log(1 - h(x^{(i)}))$$

- Repetir hasta lograr convergencia

$$\theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

Cantidad de clases

- 2 clases
 - Enfoque ya visto
- Más de 2 clases
 - Enfoque 1 vs. el resto

Ejemplo usando `fmin_bfgs` de `scipy`

Ejemplo usando fmin_bfgs de scipy

```
from numpy import loadtxt, where, zeros, e, array, log, ones, append,  
linspace  
from pylab import scatter, show, legend, xlabel, ylabel, contour, title  
from scipy.optimize import fmin_bfgs
```

```
def sigmoid(X):  
    """Compute the sigmoid function """  
    #d = zeros(shape=(X.shape))  
    den = 1.0 + e ** (-1.0 * X)  
    d = 1.0 / den  
    return d
```

Ejemplo usando fmin_bfgs de scipy

```
def cost_function_reg(theta, X, y, l):  
    """Compute the cost and partial derivatives as grads"""  
    h = sigmoid(X.dot(theta))  
    thetaR = theta[1:, 0]  
    J = (1.0 / m) * ((-y.T.dot(log(h))) - ((1 - y.T).dot(log(1.0 - h)))) \  
        + (l / (2.0 * m)) * (thetaR.T.dot(thetaR))  
  
    delta = h - y  
    sumdelta = delta.T.dot(X[:, 1])  
    grad1 = (1.0 / m) * sumdelta  
  
    XR = X[:, 1:X.shape[1]]  
    sumdelta = delta.T.dot(XR)  
  
    grad = (1.0 / m) * (sumdelta + l * thetaR)  
    out = zeros(shape=(grad.shape[0], grad.shape[1] + 1))  
  
    out[:, 0] = grad1  
    out[:, 1:] = grad  
  
    return J.flatten(), out.T.flatten()
```

Ejemplo usando fmin_bfgs de scipy

```
m, n = X.shape
```

```
y.shape = (m, 1)
```

```
it = map_feature(X[:, 0], X[:, 1])
```

```
#Initialize theta parameters
```

```
initial_theta = zeros(shape=(it.shape[1], 1))
```

```
#Set regularization parameter lambda to 1
```

```
l = 1
```

```
# Compute and display initial cost and gradient for regularized logistic  
# regression
```

```
cost, grad = cost_function_reg(initial_theta, it, y, l)
```

```
def decorated_cost(theta):
```

```
    return cost_function_reg(theta, it, y, l)
```

```
fmin_bfgs(decorated_cost, initial_theta, maxiter=400)
```

Ejemplo usando
`sklearn.linear_model.LogisticRegression`

Ejemplo usando `sklearn.linear_model.LogisticRegression`





Iris versicolor



Iris virginica



Iris setosa



Ejemplo usando sklearn.linear_model.LogisticRegression

Code source: Gaël Varoquaux

License: BSD 3 clause

import numpy **as** np

import matplotlib.pyplot **as** plt

from sklearn **import** linear_model, datasets

import some data to play with

iris = datasets.load_iris()

X = iris.data[:, :2] # we only take the first two features.

Y = iris.target

h = .02 # step size in the mesh

logreg = linear_model.LogisticRegression(C=1e5)

we create an instance of Neighbours Classifier and fit the data.

logreg.fit(X, Y)


Ejemplo usando sklearn.linear_model.LogisticRegression

```
# Plot the decision boundary. For that, we will assign a color to each
# point in the mesh [x_min, x_max][y_min, y_max].
x_min, x_max = X[:, 0].min() - .5, X[:, 0].max() + .5
y_min, y_max = X[:, 1].min() - .5, X[:, 1].max() + .5
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
Z = logreg.predict(np.c_[xx.ravel(), yy.ravel()])

# Put the result into a color plot
Z = Z.reshape(xx.shape)
plt.figure(1, figsize=(4, 3))
plt.pcolormesh(xx, yy, Z, cmap=plt.cm.Paired)
# Plot also the training points
plt.scatter(X[:, 0], X[:, 1], c=Y, edgecolors='k', cmap=plt.cm.Paired)
plt.xlabel('Sepal length')
plt.ylabel('Sepal width')

plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())
plt.xticks(())
plt.yticks(())
plt.show()
```

Ejemplo usando `sklearn.linear_model.LogisticRegression`



2. K-means

Aprendizaje automático

	<u>Unsupervised</u>	<u>Supervised</u>
<u>Continuous</u>	<ul style="list-style-type: none">• Clustering & Dimensionality Reduction<ul style="list-style-type: none">○ SVD○ PCA○ K-means	<ul style="list-style-type: none">• Regression<ul style="list-style-type: none">○ Linear○ Polynomial• Decision Trees• Random Forests
<u>Categorical</u>	<ul style="list-style-type: none">• Association Analysis<ul style="list-style-type: none">○ Apriori○ FP-Growth• Hidden Markov Model	<ul style="list-style-type: none">• Classification<ul style="list-style-type: none">○ KNN○ Trees○ Logistic Regression○ Naive-Bayes○ SVM

K-means

- Tenemos

$$X = \begin{bmatrix} \text{---} (x^{(1)})^T \text{---} \\ \text{---} (x^{(2)})^T \text{---} \\ \vdots \\ \text{---} (x^{(m)})^T \text{---} \end{bmatrix}$$

- Queremos: hacer una partición



Algoritmos de clustering basados en...

Centroides

Conectividad

Densidad


Probabilidad

*Reducción de
dimensión*


NN/Deep learning

...

Concepto de centroide



Concepto de centroide



Algoritmo k-means

- k centroides

$$\mu_1, \mu_2, \dots, \mu_k \in \mathbb{R}^n$$

- m asignaciones (una por elemento)

$$c^{(i)}$$

Algoritmo k-means

1. Inicializar los k centroides
2. Repetir hasta lograr convergencia:
 - Asignar cada elemento al centroide más cercano

For every i , set


$$c^{(i)} := \arg \min_j \|x^{(i)} - \mu_j\|^2$$

- Recalcular los centroides


For each j , set

$$\mu_j := \frac{\sum_{i=1}^m 1\{c^{(i)} = j\} x^{(i)}}{\sum_{i=1}^m 1\{c^{(i)} = j\}}$$


Algoritmo k-means, ilustrado




Algoritmo k-means, ilustrado




Algoritmo k-means, ilustrado




Algoritmo k-means, ilustrado




Algoritmo k-means, ilustrado




Algoritmo k-means, ilustrado




Algoritmo k-means, ilustrado



Algoritmo k-means, ilustrado



Algoritmo k-means, ilustrado



Función distorsión

$$J(c, \mu) = \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

- Se minimiza con **coordinate descent**. Recordar:

Repeat until convergence: {

For every i , set

$$c^{(i)} := \arg \min_j \|x^{(i)} - \mu_j\|^2$$

For each j , set

$$\mu_j := \frac{\sum_{i=1}^m 1\{c^{(i)} = j\} x^{(i)}}{\sum_{i=1}^m 1\{c^{(i)} = j\}}$$

}

- J es no convexa

Ejemplo k-means con Python

Ejemplo k-means con Python

```
import numpy as np

def cluster_points(X, mu):
    clusters = {}
    for x in X:
        bestmukey = min([(i[0], np.linalg.norm(x-mu[i[0]])) \
                        for i in enumerate(mu)], key=lambda t:t[1][0])
        try:
            clusters[bestmukey].append(x)
        except KeyError:
            clusters[bestmukey] = [x]
    return clusters


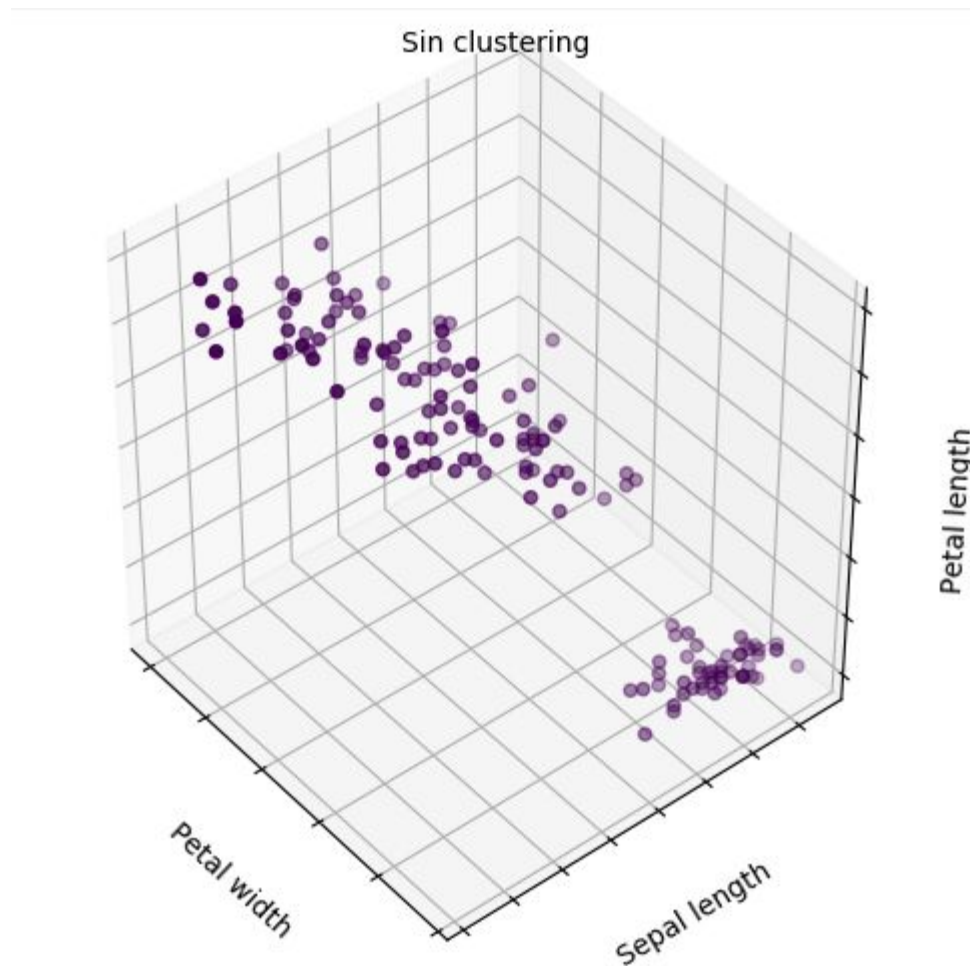
def reevaluate_centers(mu, clusters):
    newmu = []
    keys = sorted(clusters.keys())
    for k in keys:
        newmu.append(np.mean(clusters[k], axis = 0))
    return newmu
```


Ejemplo k-means con Python

```
def has_converged(mu, oldmu):  
    return (set([tuple(a) for a in mu]) == set([tuple(a) for a in oldmu]))  
  
def find_centers(X, K):  
    # Initialize to K random centers  
    oldmu = random.sample(X, K)  
    mu = random.sample(X, K)  
    while not has_converged(mu, oldmu):  
        oldmu = mu  
        # Assign all points in X to clusters  
        clusters = cluster_points(X, mu)  
        # Reevaluate centers  
        mu = reevaluate_centers(oldmu, clusters)  
    return(mu, clusters)
```

Ejemplo usando `sklearn.cluster.KMeans`

Ejemplo usando sklearn.cluster.KMeans



Ejemplo usando sklearn.cluster.KMeans

```
import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
```

```
from sklearn.cluster import KMeans
from sklearn import datasets
```

```
np.random.seed(5)
```

```
centers = [[1, 1], [-1, -1], [1, -1]]
iris = datasets.load_iris()
X = iris.data
y = iris.target
```

```
estimators = {'k_means_iris_2': KMeans(n_clusters=2),
              'k_means_iris_3': KMeans(n_clusters=3),
              'k_means_iris_5': KMeans(n_clusters=5)}
```

Ejemplo usando sklearn.cluster.KMeans

```
fignum = 1
for name, est in estimators.items():
    fig = plt.figure(fignum, figsize=(4, 3))
    plt.clf()
    #add title
    fig.suptitle(name, fontsize=10)
    ax = Axes3D(fig, rect=[0, 0, .95, 1], elev=48, azim=134)

    plt.cla()
    est.fit(X)
    labels = est.labels_

    ax.scatter(X[:, 3], X[:, 0], X[:, 2], c=labels.astype(np.float))

    ax.set_xlabel('Petal width')
    ax.set_ylabel('Sepal length')
    ax.set_zlabel('Petal length')
    fignum = fignum + 1
```

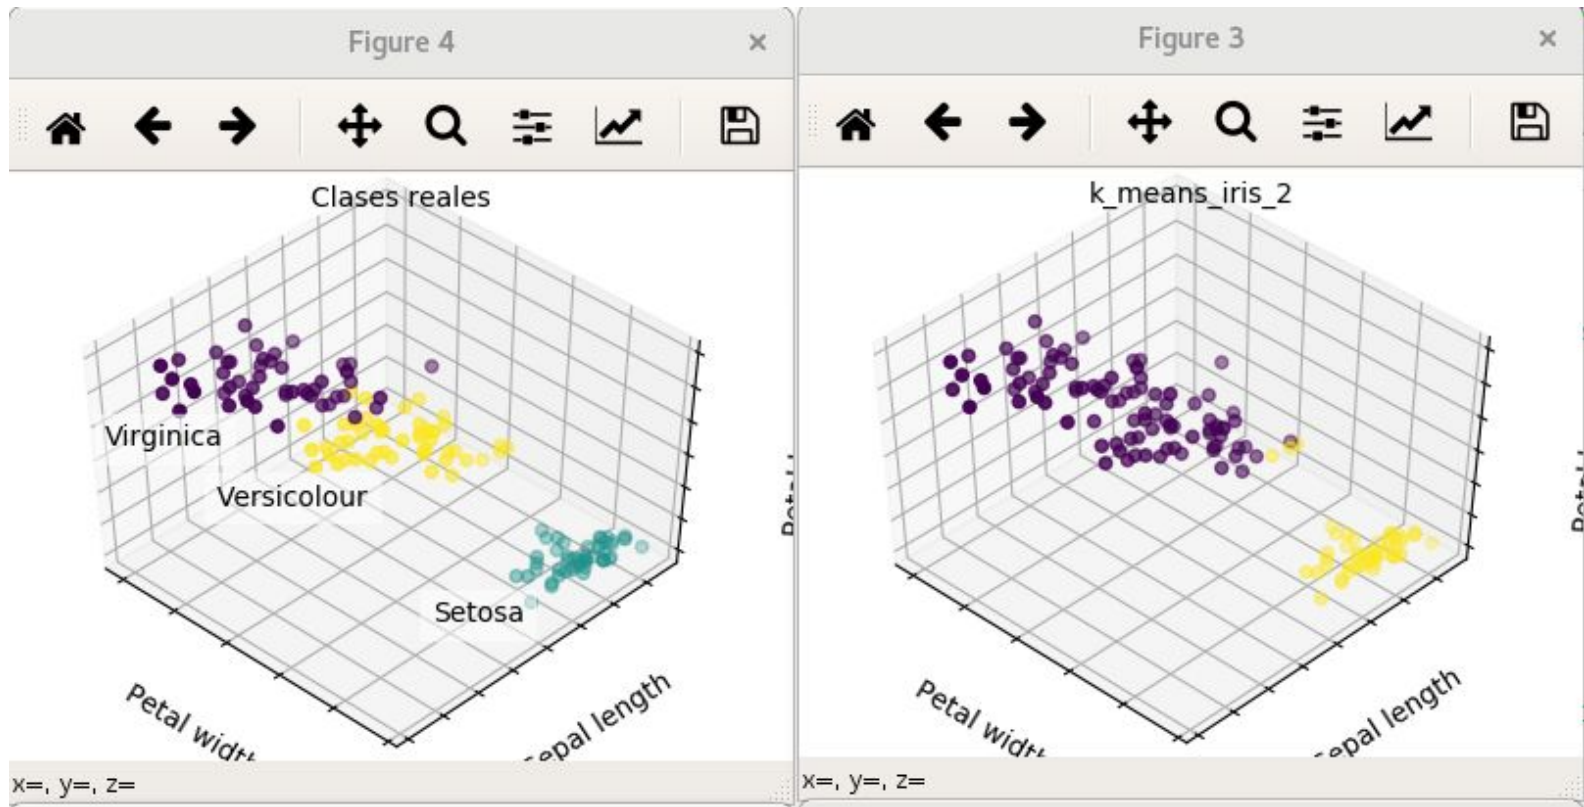
Ejemplo usando sklearn.cluster.KMeans

Plot the ground truth

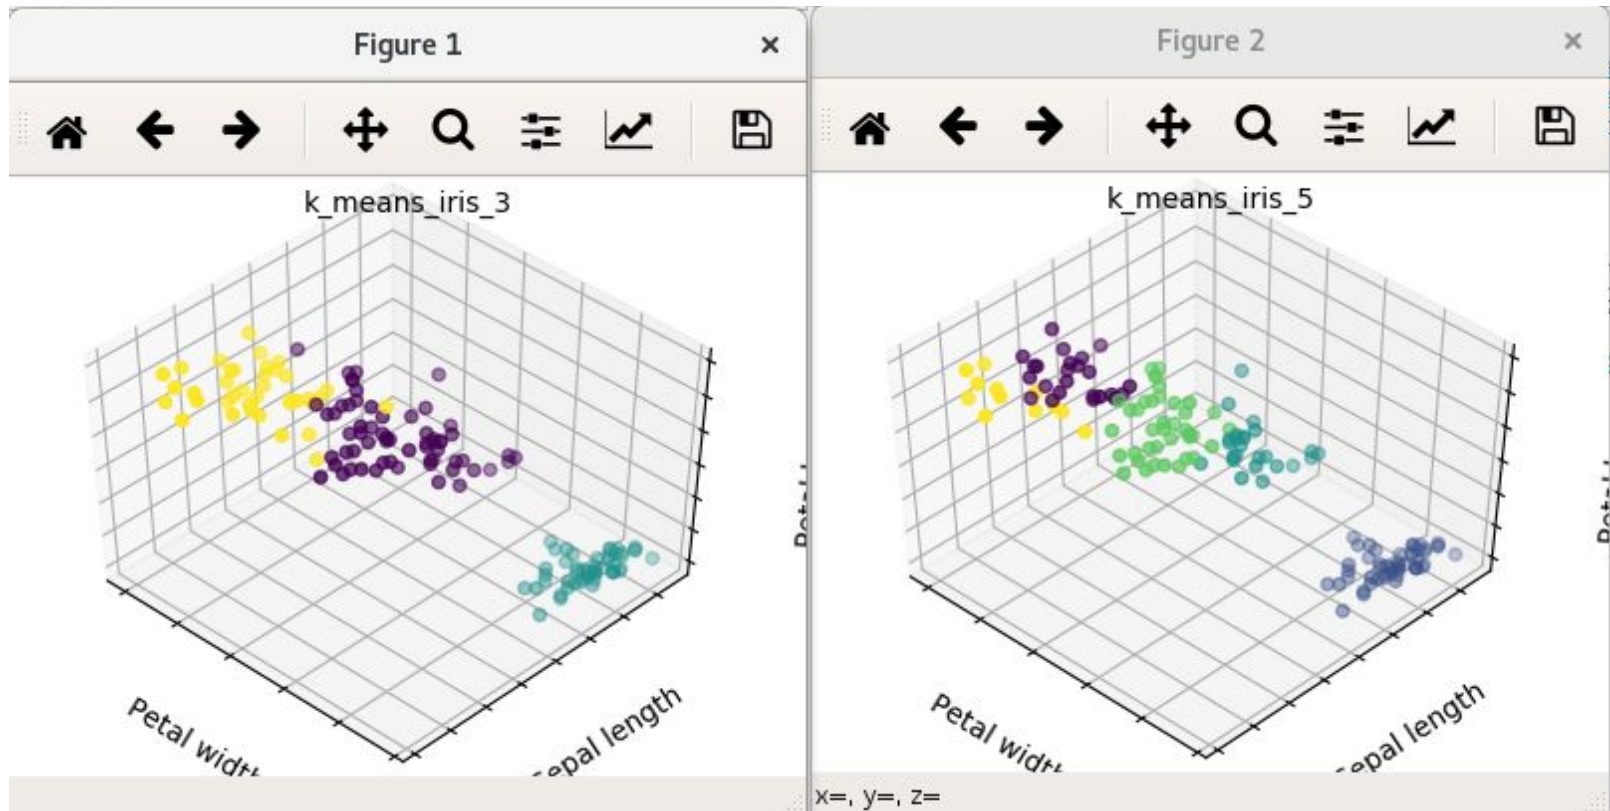
```
fig = plt.figure(figsize=(4, 3))
plt.clf()
ax = Axes3D(fig, rect=[0, 0, .95, 1], elev=48, azim=134)
fig.suptitle('Clases reales', fontsize=10)
plt.cla()
```

```
for name, label in [('Setosa', 0),
                    ('Versicolour', 1),
                    ('Virginica', 2)]:
    ax.text3D(X[y == label, 3].mean(),
              X[y == label, 0].mean() + 1.5,
              X[y == label, 2].mean(), name,
              horizontalalignment='center',
              bbox=dict(alpha=.5, edgecolor='w', facecolor='w'))
# Reorder the labels to have colors matching the cluster results
y = np.choose(y, [1, 2, 0]).astype(np.float)
ax.scatter(X[:, 3], X[:, 0], X[:, 2], c=y)
```

Ejemplo usando sklearn.cluster.KMeans



Ejemplo usando sklearn.cluster.KMeans



Fuentes consultadas

- CS229 Supervised Learning (<http://cs229.stanford.edu/notes/cs229-notes1.pdf>)
- Slides de Universität Leipzig, Classification using Logistic Regression
- Documentación de scikit LogisticRegression
http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html
- Logistic Regression 3-class classifier
(http://scikit-learn.org/stable/auto_examples/linear_model/plot_iris_logistic.html#sphx-glr-auto-examples-linear-model-plot-iris-logistic-py)
- Machine Learning with Python - Logistic Regression
(<http://aimotion.blogspot.mx/2011/11/machine-learning-with-python-logistic.html>)
- Logistic Regression (<https://github.com/perborgen/LogisticRegression/blob/master/logistic.py>)
- CS229 Lecture notes, The k-means clustering algorithm
(<http://cs229.stanford.edu/notes/cs229-notes7a.pdf>)
- Clustering With K-Means in Python
(<https://datasciencelab.wordpress.com/2013/12/12/clustering-with-k-means-in-python/>)
- K-means clustering
(http://scikit-learn.org/stable/auto_examples/cluster/plot_cluster_iris.html#sphx-glr-auto-examples-cluster-plot-cluster-iris-py)

1. Regresión logística

2. K-means



Meetup Monterrey Data Science & Engineering



CoWo, Monterrey, Nuevo León, México



22 de junio de 2017



Rafael Rodríguez Morales (rafarodrz@gmail.com)